

# Potato Disease Classification

Dataset credits: <https://www.kaggle.com/arjuntejaswi/plant-village>

## Importing dependencies

```
In [1]: import tensorflow as tf
        from tensorflow.keras import models, layers
        import matplotlib.pyplot as plt
```

## Set all the Constants

```
In [2]: IMAGE_SIZE = 256
        BATCH_SIZE = 32
        CHANNELS = 3
        EPOCHS = 10
```

## Import data into tensorflow dataset object

We will use `image_dataset_from_directory` api to load all images in tensorflow dataset:

[https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image\\_dataset\\_from\\_directory](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image_dataset_from_directory)

```
In [3]: dataset = tf.keras.preprocessing.image_dataset_from_directory(
        "/content/drive/MyDrive/PlantVillage",
        shuffle = True,
        image_size = (IMAGE_SIZE, IMAGE_SIZE),
        batch_size = BATCH_SIZE
    )
```

Found 2152 files belonging to 3 classes.

```
In [4]: class_names = dataset.class_names
class_names
```

```
Out[4]: ['Potato__Early_blight', 'Potato__Late_blight', 'Potato__healthy']
```

```
In [7]: for image_batch, labels_batch in dataset.take(1):
        print(image_batch.shape)
        print(labels_batch.numpy())
```

```
(32, 256, 256, 3)
[2 0 1 1 1 0 1 0 1 1 1 0 1 0 0 0 0 1 0 0 1 1 1 1 1 1 1 0 1 1 1 0]
```

As you can see above, each element in the dataset is a tuple. First element is a batch of 32 elements of images. Second element is a batch of 32 elements of class labels

### Visualize some of the images from our dataset

```
In [8]: plt.figure(figsize=(10,10))
for image_batch, label_batch in dataset.take(1):
    for i in range(12):
        ax = plt.subplot(3,4,1+i)
        plt.imshow(image_batch[i].numpy().astype("uint8"))
        plt.title(class_names[label_batch[i]])
        plt.axis("off")
```

Potato\_\_Early\_blight



Potato\_\_Late\_blight



Potato\_\_Late\_blight



Potato\_\_Late\_blight



Potato\_\_Early\_blight



Potato\_\_Late\_blight



Potato\_\_Late\_blight



Potato\_\_Late\_blight



Potato\_\_Late\_blight



Potato\_\_healthy



Potato\_\_Late\_blight



Potato\_\_Early\_blight



### Function to Split Dataset

Dataset should be bifurcated into 3 subsets, namely:

1. Training: Dataset to be used while training.
2. Validation: Dataset to be tested against while training
3. Test: Dataset to be tested against after we trained a model

```
In [9]: len(dataset)
```

```
Out[9]: 68
```

```
In [10]: train_size = 0.8  
len(dataset)*train_size
```

```
Out[10]: 54.400000000000006
```

```
In [11]: train_ds = dataset.take(54)  
len(train_ds)
```

```
Out[11]: 54
```

```
In [12]: test_ds = dataset.skip(54)  
len(test_ds)
```

```
Out[12]: 14
```

```
In [13]: val_size = 0.1  
len(dataset)*val_size
```

```
Out[13]: 6.800000000000001
```

```
In [14]: val_ds = test_ds.take(6)  
len(val_ds)
```

```
Out[14]: 6
```

```
In [15]: test_ds = test_ds.skip(6)  
len(test_ds)
```

```
Out[15]: 8
```

```
In [16]: def get_dataset_partitions_tf(ds, train_split=0.8, val_split=0.1, test_split=0.1, shuffle=True, shuffle_size=10000):
    assert (train_split + test_split + val_split) == 1

    ds_size = len(ds)

    if shuffle:
        ds = ds.shuffle(shuffle_size, seed=12)

    train_size = int(train_split * ds_size)
    val_size = int(val_split * ds_size)

    val_size = int(val_split * ds_size)

    train_ds = ds.take(train_size)
    val_ds = ds.skip(train_size).take(val_size)
    test_ds = ds.skip(train_size).skip(val_size)

    return train_ds, val_ds, test_ds
```

```
In [17]: train_ds, val_ds, test_ds = get_dataset_partitions_tf(dataset)
```

```
In [18]: len(train_ds)
```

```
Out[18]: 54
```

```
In [19]: len(val_ds)
```

```
Out[19]: 6
```

```
In [20]: len(test_ds)
```

```
Out[20]: 8
```

### Cache, Shuffle, and Prefetch the Dataset

```
In [21]: train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size= tf.data.AUTOTUNE)
val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size= tf.data.AUTOTUNE)
test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size= tf.data.AUTOTUNE)
```

## Building the Model

### Creating a Layer for Resizing and Normalization

Before we feed our images to network, we should be resizing it to the desired size. Moreover, to improve model performance, we should normalize the image pixel value (keeping them in range 0 and 1 by dividing by 255). This should happen while training as well as inference. Hence we can add that as a layer in our Sequential Model.

```
In [22]: resize_and_rescale = tf.keras.Sequential([
    layers.experimental.preprocessing.Resizing(IMAGE_SIZE, IMAGE_SIZE),
    layers.experimental.preprocessing.Rescaling(1.0/255),
])
```

## Data Augmentation

Data Augmentation is needed when we have less data, this boosts the accuracy of our model by augmenting the data.

```
In [23]: data_augmentation = tf.keras.Sequential([
    layers.experimental.preprocessing.RandomFlip("horizontal_and_vertical"),
    layers.experimental.preprocessing.RandomRotation(0.2),
])
```

---

## Applying Data Augmentation to Train Dataset

```
In [24]: train_ds = train_ds.map(  
         lambda x, y: (data_augmentation(x, training=True), y)  
         ).prefetch(buffer_size=tf.data.AUTOTUNE)
```

## Model Architecture

We use a CNN coupled with a Softmax activation in the output layer. We also add the initial layers for resizing, normalization and Data Augmentation.

We are going to use convolutional neural network (CNN) here. CNN is popular for image classification tasks. Watch below video to understand fundamentals of CNN

```
In [25]: input_shape = (BATCH_SIZE, IMAGE_SIZE, IMAGE_SIZE, CHANNELS)  
         n_classes = 3  
  
         model = models.Sequential([  
             resize_and_rescale,  
             layers.Conv2D(32, kernel_size = (3,3), activation='relu', input_shape=input_shape),  
             layers.MaxPooling2D((2, 2)),  
             layers.Conv2D(64, kernel_size = (3,3), activation='relu'),  
             layers.MaxPooling2D((2, 2)),  
             layers.Conv2D(64, kernel_size = (3,3), activation='relu'),  
             layers.MaxPooling2D((2, 2)),  
             layers.Conv2D(64, (3, 3), activation='relu'),  
             layers.MaxPooling2D((2, 2)),  
             layers.Conv2D(64, (3, 3), activation='relu'),  
             layers.MaxPooling2D((2, 2)),  
             layers.Conv2D(64, (3, 3), activation='relu'),  
             layers.MaxPooling2D((2, 2)),  
             layers.Flatten(),
```

```

        layers.Dense(64, activation='relu'),
        layers.Dense(n_classes, activation='softmax'),
    ])

    model.build(input_shape=input_shape)

```

In [26]: `model.summary()`

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
sequential (Sequential)	(32, 256, 256, 3)	0
conv2d (Conv2D)	(32, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(32, 127, 127, 32)	0
conv2d_1 (Conv2D)	(32, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(32, 62, 62, 64)	0
conv2d_2 (Conv2D)	(32, 60, 60, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(32, 30, 30, 64)	0
conv2d_3 (Conv2D)	(32, 28, 28, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(32, 14, 14, 64)	0
conv2d_4 (Conv2D)	(32, 12, 12, 64)	36928



conv2d_4 (Conv2D)	(32, 12, 12, 64)	36928
max_pooling2d_4 (MaxPooling 2D)	(32, 6, 6, 64)	0
conv2d_5 (Conv2D)	(32, 4, 4, 64)	36928
max_pooling2d_5 (MaxPooling 2D)	(32, 2, 2, 64)	0
flatten (Flatten)	(32, 256)	0
dense (Dense)	(32, 64)	16448
dense_1 (Dense)	(32, 3)	195

```
=====
Total params: 183,747
Trainable params: 183,747
Non-trainable params: 0
```

---

## Compiling the Model

We use adam Optimizer, SparseCategoricalCrossentropy for losses, accuracy as a metric

```
In [27]: model.compile(
    optimizer = 'adam',
    loss = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics = ['accuracy']
)
```

```
In [28]: history = model.fit(
    train_ds,
    epochs = EPOCHS,
    batch_size = BATCH_SIZE,
    verbose = 1,
    validation_data = val_ds
)
```

Epoch 1/10

54/54 [=====] - 397s 3s/step - loss: 0.9037 - accuracy: 0.5023 - val\_loss: 0.8763 - val\_accuracy: 0.6406

Epoch 2/10

54/54 [=====] - 175s 3s/step - loss: 0.6836 - accuracy: 0.7274 - val\_loss: 1.1239 - val\_accuracy: 0.5365

Epoch 3/10

54/54 [=====] - 174s 3s/step - loss: 0.4899 - accuracy: 0.7998 - val\_loss: 0.5737 - val\_accuracy: 0.7448

Epoch 4/10

54/54 [=====] - 174s 3s/step - loss: 0.3376 - accuracy: 0.8611 - val\_loss: 0.2391 - val\_accuracy: 0.9062

Epoch 5/10

54/54 [=====] - 173s 3s/step - loss: 0.2831 - accuracy: 0.8866 - val\_loss: 0.4325 - val\_accuracy: 0.7917

Epoch 6/10

54/54 [=====] - 173s 3s/step - loss: 0.2391 - accuracy: 0.9057 - val\_loss: 0.1906 - val\_accuracy: 0.9115

Epoch 7/10

54/54 [=====] - 173s 3s/step - loss: 0.2075 - accuracy: 0.9201 - val\_loss: 0.2772 - val\_accuracy: 0.9062

Epoch 8/10

54/54 [=====] - 173s 3s/step - loss: 0.2129 - accuracy: 0.9161 - val\_loss: 0.1976 - val\_accuracy: 0.9271

Epoch 9/10

54/54 [=====] - 174s 3s/step - loss: 0.1979 - accuracy: 0.9144 - val\_loss: 0.2772 - val\_accuracy: 0.8802

Epoch 10/10

54/54 [=====] - 173s 3s/step - loss: 0.2155 - accuracy: 0.9109 - val\_loss: 0.1297 - val\_accuracy: 0.9427

```
In [30]: scores = model.evaluate(test_ds)
```

```
8/8 [=====] - 11s 762ms/step - loss: 0.1287 - accuracy: 0.9570
```

**You can see above that we get 95.70% accuracy for our test dataset. This is considered to be a pretty good accuracy**

```
In [31]: scores
```

```
Out[31]: [0.1286834180355072, 0.95703125]
```

**Scores is just a list containing loss and accuracy value**

```
In [32]: history.params
```

```
Out[32]: {'epochs': 10, 'steps': 54, 'verbose': 1}
```

```
In [33]: history.history.keys()
```

```
Out[33]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

loss, accuracy, val loss etc are a python list containing values of loss, accuracy etc at the end of each epoch

```
In [34]: type(history.history['loss'])
```

```
Out[34]: list
```

```
In [35]: len(history.history['loss'])
```

```
Out[35]: 10
```


```
In [36]: history.history['loss'][:5]
```

```
Out[36]: [0.9037485718727112,  
          0.683612048625946,  
          0.48990094661712646,  
          0.33762332797050476,  
          0.28306978940963745]
```

```
In [37]: acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']  
  
loss = history.history['loss']  
val_loss = history.history['val_loss']
```

## Run prediction on a sample image

```
In [38]: plt.figure(figsize=(8,8))  
plt.subplot(1, 2, 1)  
plt.plot(range(EPOCHS), acc, label = 'Training Accuracy')  
plt.plot(range(EPOCHS), val_acc, label = 'Validation Loss')  
plt.legend(loc = 'lower right')  
plt.title('Training and Validation Accuracy')  
  
plt.subplot(1, 2, 2)  
plt.plot(range(EPOCHS), loss, label='Training Loss')  
plt.plot(range(EPOCHS), val_loss, label='Validation Loss')  
plt.legend(loc='upper right')  
plt.title('Training and Validation Loss')  
plt.show()
```



```
In [39]: import numpy as np
for images_batch, labels_batch in test_ds.take(1):


    first_image = images_batch[0].numpy().astype('uint8')
    first_label = labels_batch[0].numpy()
```

```
first_image = images_batch[0].numpy().astype('uint8')
first_label = labels_batch[0].numpy()

print("first image to predict")
plt.imshow(first_image)
print("actual label:", class_names[first_label])

batch_prediction = model.predict(images_batch)
print("predicted label:", class_names[np.argmax(batch_prediction[0])])
```

first image to predict  
actual label: Potato\_\_Late\_blight  
predicted label: Potato\_\_Late\_blight



## Write a function for inference

```
In [40]: def predict(model, img):
img_array = tf.keras.preprocessing.image.img_to_array(images[i].numpy())
img_array = tf.expand_dims(img_array, 0)

predictions = model.predict(img_array)

predicted_class = class_names[np.argmax(predictions[0])]
confidence = round(100 * (np.max(predictions[0])), 2)
return predicted_class, confidence
```

## Now run inference on few sample images

```
In [41]: plt.figure(figsize=(15, 15))
for images, labels in test_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))

        predicted_class, confidence = predict(model, images[i].numpy())
        actual_class = class_names[labels[i]]

        plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class}.\n Confidence: {confidence}%")

        plt.axis("off")
```

Actual: Potato\_\_healthy,  
Predicted: Potato\_\_healthy.  
Confidence: 95.59%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Early\_blight.  
Confidence: 83.0%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 91.55%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 98.09%



Actual: Potato\_\_Early\_blight,  
Predicted: Potato\_\_Early\_blight.  
Confidence: 97.4%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 98.71%



Actual: Potato\_\_Early\_blight,  
Predicted: Potato\_\_Early\_blight.  
Confidence: 100.0%



Actual: Potato\_\_Early\_blight,  
Predicted: Potato\_\_Early\_blight.  
Confidence: 96.35%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 96.61%

